

OM performance test result (prototype)

Scope:

- Leader side execution
- Optimized flow for OBS create/commit (and involves overwrite of previous key)

Optimization key points:

1. Parallel execution of different keys (Granular lock for key using stripped lock)
2. No Caching
3. Batching of ratis db update to flush to all nodes
4. Flow optimization removing redundant and un-necessary operation

Environment:

Cluster with 16 DNs, OM 3 node HA, SCM 3 node HA

Master nodes:

CPU	2 x Intel(R) Xeon(R) Gold 6230 CPU @ 2.10GHz/20 cores
memory	384GB (12 x 32GB DDR4 @ 2933MHz)
OS Boot	Cisco Boot optimized M.2 Raid controller with 2 x 240GB SATA SSD
SSD	3.8TB SATA SSD Enterprise Value
Storage Controller	Cisco 12G Modular Raid Controller with 2GB cache
Network Adapter	Cisco UCS VIC 1387 2 x 40Gbps ports x8 PCIe Gen3

Datanodes:

CPU	2 x Intel(R) Xeon(R) Gold 6262V CPU @ 1.90GHz/24 cores
memory	384GB (12 x 32GB DDR4 @ 2933MHz)
OS Boot	Cisco Boot optimized M.2 Raid controller with 2 x 240GB SATA SSD
NVMe	10 x 8TB Intel P4510 U.2 High Performance Value
Network Adapter	Cisco UCS VIC 1387 2 x 40Gbps ports x8 PCIe Gen3

Comparison:

sno	Category	Old Flow	New Flow
1	Operation / Second (key create / commit)	12K+	40K+
2	Key Commit / Second	5.9K+	20K+ (3.3 times)
3	CPU Utilization Leader	16% (unable to increase load)	33%
4	CPU Utilization Follower	6% above	4% below
5			

Further Improvement

Further improvements can increase to **30K/Second**:

1. Caching of open Key for commit to avoid get from table
2. Codec Buffer encoding/decoding to be improved as per new flow
3. GetBlock caching at OM

Note: This is from the fact,

- Mocking getBlock and ranger reached capability to 40k+ key commit / second (almost double)

Few more area of improvement to increase caps:

1. Ranger Validation - to optimize / cached frequent call
2. Ratis upgrade to new version and further improvement

Command used for test

(thread in increased and multiple instance run to increase load under various test):

```
ozone freon ockrw -n 10000000 -t 100 --percentage-read 0 --size 0 -r 1000000 -v voltest -b buckettest -p performanceTest
```

New Flow:

This is data with one freon running. When load increases with multiple freon, its crosses 21k easily.

key-read-write-list

count = 10000000

mean rate = 19565.78 calls/second

1-minute rate = 19989.40 calls/second

5-minute rate = 18552.83 calls/second

15-minute rate = 16043.74 calls/second

min = 3.02 milliseconds

max = 25.84 milliseconds

mean = 4.67 milliseconds

stddev = 2.02 milliseconds

median = 4.26 milliseconds

75% <= 4.64 milliseconds

95% <= 6.40 milliseconds

98% <= 14.51 milliseconds

99% <= 16.30 milliseconds

99.9% <= 18.39 milliseconds

Total execution time (sec): 513

Failures: 0

Successful executions: 10000000






Old flow:

key-read-write-list


count = 10000000
 mean rate = 5913.05 calls/second
 1-minute rate = 5822.04 calls/second
 5-minute rate = 5859.77 calls/second
 15-minute rate = 5567.46 calls/second
 min = 12.08 milliseconds
 max = 927.93 milliseconds
 mean = 17.38 milliseconds
 stddev = 27.01 milliseconds
 median = 14.94 milliseconds
 75% <= 16.25 milliseconds
 95% <= 32.03 milliseconds
 98% <= 35.81 milliseconds
 99% <= 39.04 milliseconds
 99.9% <= 76.52 milliseconds

Total execution time (sec): 1693 ⇒ **28 minute 29 second**

GC Time - jvm_metrics_gc_time_millis (OM)



GC Count - jvm_metrics_gc_count (OM)



~ OM CPU Load Metrics

Cpu Load - jvm_metrics_cpu_jvm_load (OM)

